

# California Housing Price Prediction Machine Learning Project

September 26, 2022

```
[1]: #Import Necessary Libraries:
import pandas as pd
import numpy as np

from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.tree import DecisionTreeRegressor
import statsmodels.formula.api as smf

from sklearn.metrics import mean_squared_error, r2_score
from math import sqrt

import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

from matplotlib.axes._axes import _log as matplotlib_axes_logger
matplotlib_axes_logger.setLevel('ERROR')
```

```
[2]: df_house=pd.read_excel("1553768847_housing.xlsx")
df_house.head()
```

```
[2]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41	880	129.0	
1	-122.22	37.86	21	7099	1106.0	
2	-122.24	37.85	52	1467	190.0	
3	-122.25	37.85	52	1274	235.0	
4	-122.25	37.85	52	1627	280.0	

	population	households	median_income	ocean_proximity	median_house_value
0	322	126	8.3252	NEAR BAY	452600
1	2401	1138	8.3014	NEAR BAY	358500
2	496	177	7.2574	NEAR BAY	352100
3	558	219	5.6431	NEAR BAY	341300

4	565	259	3.8462	NEAR BAY	342200
---	-----	-----	--------	----------	--------

```
[3]: import math
      print(math.log(452600))
```

13.022764012181574

```
[4]: df_house.columns
```

```
[4]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
          'total_bedrooms', 'population', 'households', 'median_income',
          'ocean_proximity', 'median_house_value'],
          dtype='object')
```

```
[5]: df_house.isnull().sum()
```

```
[5]: longitude          0
     latitude          0
     housing_median_age  0
     total_rooms        0
     total_bedrooms    207
     population         0
     households         0
     median_income      0
     ocean_proximity    0
     median_house_value  0
     dtype: int64
```

```
[6]: df_house.total_bedrooms=df_house.total_bedrooms.fillna(df_house.total_bedrooms.
     ↪mean())
     df_house.isnull().sum()
```

```
[6]: longitude          0
     latitude          0
     housing_median_age  0
     total_rooms        0
     total_bedrooms    0
     population         0
     households         0
     median_income      0
     ocean_proximity    0
     median_house_value  0
     dtype: int64
```

```
[7]: le = LabelEncoder()
     df_house['ocean_proximity']=le.fit_transform(df_house['ocean_proximity'])
```

```
[8]: # Get column names first
names = df_house.columns
# Create the Scaler object
scaler = StandardScaler()
# Fit your data on the scaler object
scaled_df = scaler.fit_transform(df_house)
scaled_df = pd.DataFrame(scaled_df, columns=names)
scaled_df.head()
```

```
[8]: longitude latitude housing_median_age total_rooms total_bedrooms \
0 -1.327835 1.052548 0.982143 -0.804819 -0.975228
1 -1.322844 1.043185 -0.607019 2.045890 1.355088
2 -1.332827 1.038503 1.856182 -0.535746 -0.829732
3 -1.337818 1.038503 1.856182 -0.624215 -0.722399
4 -1.337818 1.038503 1.856182 -0.462404 -0.615066

population households median_income ocean_proximity median_house_value
0 -0.974429 -0.977033 2.344766 1.291089 2.129631
1 0.861439 1.669961 2.332238 1.291089 1.314156
2 -0.820777 -0.843637 1.782699 1.291089 1.258693
3 -0.766028 -0.733781 0.932968 1.291089 1.165100
4 -0.759847 -0.629157 -0.012881 1.291089 1.172900
```

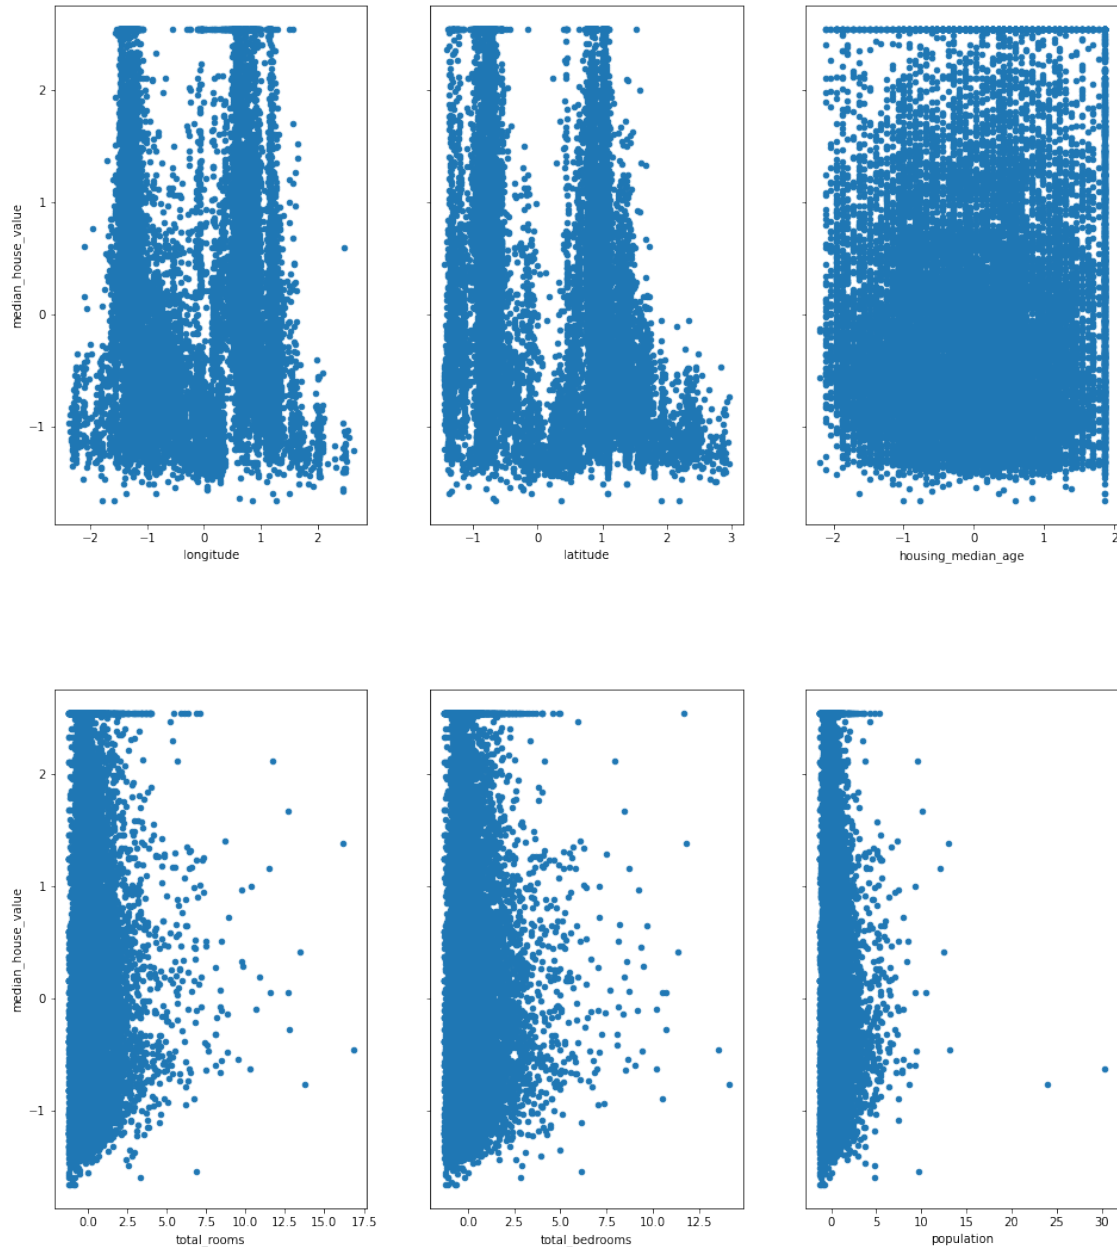
```
[9]: #plot graphs
fig,axs=plt.subplots(1,3,sharey=True)
scaled_df.
    ↳plot(kind='scatter',x='longitude',y='median_house_value',ax=axs[0],figsize=(16,8))
scaled_df.
    ↳plot(kind='scatter',x='latitude',y='median_house_value',ax=axs[1],figsize=(16,8))
scaled_df.
    ↳plot(kind='scatter',x='housing_median_age',y='median_house_value',ax=axs[2],figsize=(16,8))


#plot graphs
fig,axs=plt.subplots(1,3,sharey=True)
scaled_df.
    ↳plot(kind='scatter',x='total_rooms',y='median_house_value',ax=axs[0],figsize=(16,8))
scaled_df.
    ↳plot(kind='scatter',x='total_bedrooms',y='median_house_value',ax=axs[1],figsize=(16,8))
scaled_df.
    ↳plot(kind='scatter',x='population',y='median_house_value',ax=axs[2],figsize=(16,8))

#plot graphs
fig,axs=plt.subplots(1,3,sharey=True)
scaled_df.
    ↳plot(kind='scatter',x='households',y='median_house_value',ax=axs[0],figsize=(16,8))
```


```
scaled_df.  
    ↪ plot(kind='scatter',x='median_income',y='median_house_value',ax=axis[1],figsize=(16,8))  
scaled_df.  
    ↪ plot(kind='scatter',x='ocean_proximity',y='median_house_value',ax=axis[2],figsize=(16,8))
```



[9]: <AxesSubplot:xlabel='ocean\_proximity', ylabel='median\_house\_value'>









```
[10]: for column in scaled_df:
    plt.figure()
    sns.boxplot(x=scaled_df[column])
```

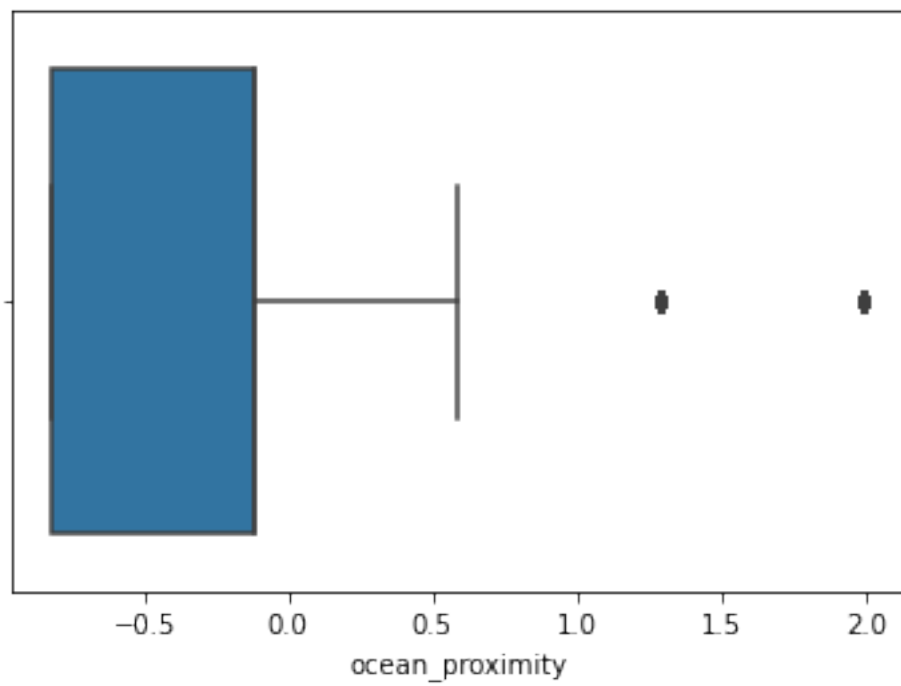
















```
[11]: X_Features=['longitude', 'latitude', 'housing_median_age', 'total_rooms',
               'total_bedrooms', 'population', 'households', 'median_income',
               'ocean_proximity']
```

```
X=scaled_df[X_Features]
```

```
Y=scaled_df['median_house_value']
```

```
print(type(X))
```

```
print(type(Y))
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
<class 'pandas.core.series.Series'>
```

```
[12]: print(df_house.shape)
```

```
print(X.shape)
```

```
print(Y.shape)
```

```
(20640, 10)
```

```
(20640, 9)
```

```
(20640,)
```

```
[13]: from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.2,random_state=1)
```

```
print (x_train.shape, y_train.shape)
```

```
print (x_test.shape, y_test.shape)
```

```
(16512, 9) (16512,)
(4128, 9) (4128,)
```

```
[14]: linreg=LinearRegression()
linreg.fit(x_train,y_train)
```

```
[14]: LinearRegression()
```

```
[15]: y_predict = linreg.predict(x_test)
```

```
[16]: print(sqrt(mean_squared_error(y_test,y_predict)))
print((r2_score(y_test,y_predict)))
```

```
0.6056598120301221
0.6276223517950295
```

```
[17]: dtreg=DecisionTreeRegressor()
dtreg.fit(x_train,y_train)
```

```
[17]: DecisionTreeRegressor()
```

```
[18]: y_predict = dtreg.predict(x_test)
print(sqrt(mean_squared_error(y_test,y_predict)))
print((r2_score(y_test,y_predict)))
```

```
0.5925633790037315
0.6435523907257298
```

```
[21]: lassoreg=Lasso(alpha=0.001,normalize=True)
lassoreg.fit(x_train,y_train)
print(sqrt(mean_squared_error(y_test,lassoreg.predict(x_test))))
print('R2 Value/Coefficient of determination:{}'.format(lassoreg.
↪score(x_test,y_test)))
```

```
0.719314096707071
R2 Value/Coefficient of determination:0.4747534206169961
```

```
[22]: ridgereg=Ridge(alpha=0.001,normalize=True)
ridgereg.fit(x_train,y_train)
print(sqrt(mean_squared_error(y_test,ridgereg.predict(x_test))))
print('R2 Value/Coefficient of determination:{}'.format(ridgereg.
↪score(x_test,y_test)))
```

```
0.6056048844852343
R2 Value/Coefficient of determination:0.6276898909055972
```

```
[23]: from sklearn.linear_model import ElasticNet
elasticreg=ElasticNet(alpha=0.001,normalize=True)
```

```

elasticreg.fit(x_train,y_train)
print(sqrt(mean_squared_error(y_test,elasticreg.predict(x_test))))
print('R2 Value/Coefficient of determination:{}'.format(elasticreg.
↪score(x_test,y_test)))

```

0.944358169398106

R2 Value/Coefficient of determination:0.09468529806704551

```

[24]: lm=smf.ols(formula='median_house_value ~_
↪longitude+latitude+housing_median_age+total_rooms+total_bedrooms+population+households+medi
↪fit()

```

```

[25]: lm.summary()

```

```

[25]: <class 'statsmodels.iolib.summary.Summary'>
"""

```

```

                                OLS Regression Results
=====
Dep. Variable:    median_house_value    R-squared:                0.636
Model:                            OLS    Adj. R-squared:           0.635
Method:                    Least Squares    F-statistic:            3999.
Date:                Mon, 26 Sep 2022    Prob (F-statistic):       0.00
Time:                  15:17:42    Log-Likelihood:        -18868.
No. Observations:        20640    AIC:                   3.776e+04
Df Residuals:            20630    BIC:                   3.783e+04
Df Model:                 9
Covariance Type:            nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
Intercept          -3.469e-17      0.004  -8.26e-15      1.000      -0.008
0.008
longitude          -0.7393      0.013   -57.263      0.000      -0.765
-0.714
latitude           -0.7858      0.013   -61.664      0.000      -0.811
-0.761
housing_median_age    0.1248      0.005    26.447      0.000      0.116
0.134
total_rooms         -0.1265      0.015    -8.609      0.000      -0.155
-0.098
total_bedrooms       0.2995      0.022    13.630      0.000      0.256
0.343
population          -0.3907      0.011   -36.927      0.000      -0.411
-0.370

```

households	0.2589	0.022	11.515	0.000	0.215
0.303					
median_income	0.6549	0.005	119.287	0.000	0.644
0.666					
ocean_proximity	0.0009	0.005	0.190	0.850	-0.008
0.010					

```
=====
Omnibus:                    5037.491    Durbin-Watson:                    0.965
Prob(Omnibus):              0.000    Jarque-Bera (JB):              18953.000
Skew:                      1.184    Prob(JB):                      0.00
Kurtosis:                  7.054    Cond. No.                      14.2
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
"""
```

```
[26]: x_train_Income=x_train[['median_income']]
      x_test_Income=x_test[['median_income']]
      print(x_train_Income.shape)
      print(y_train.shape)
```

```
(16512, 1)
(16512,)
```


```
[27]: linreg=LinearRegression()
      linreg.fit(x_train_Income,y_train)
      y_predict = linreg.predict(x_test_Income)
      #print intercept and coefficient of the linear equation
      print(linreg.intercept_, linreg.coef_)
      print(sqrt(mean_squared_error(y_test,y_predict)))
      print((r2_score(y_test,y_predict)))
```

```
0.005623019866893164 [0.69238221]
0.7212595914243148
0.47190835934467734
```

```
[28]: #plot least square line
      scaled_df.plot(kind='scatter',x='median_income',y='median_house_value')
      plt.plot(x_test_Income,y_predict,c='red',linewidth=2)
```

```
[28]: [

```



```
[29]: lm=smf.ols(formula='median_house_value ~ median_income',data=scaled_df).fit()
lm.summary()
```

```
[29]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:    median_house_value    R-squared:                0.473
Model:                            OLS    Adj. R-squared:           0.473
Method:                    Least Squares    F-statistic:            1.856e+04
Date:                Mon, 26 Sep 2022    Prob (F-statistic):       0.00
Time:                  15:18:49    Log-Likelihood:         -22668.
No. Observations:        20640    AIC:                   4.534e+04
Df Residuals:            20638    BIC:                   4.536e+04
Df Model:                 1
Covariance Type:        nonrobust
=====
=
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-
Intercept      1.735e-16      0.005      3.43e-14      1.000      -0.010
0.010
```

```

median_income      0.6881      0.005      136.223      0.000      0.678
0.698
=====
Omnibus:                4245.795      Durbin-Watson:                0.655
Prob(Omnibus):           0.000      Jarque-Bera (JB):            9273.446
Skew:                    1.191      Prob(JB):                    0.00
Kurtosis:                5.260      Cond. No.                    1.00
=====

```

```

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""

```

```
[ ]:
```